

Applied Deep Learning for NLP


Week 3 - Text Representations

Juan Carlos Medina Serrano

Technische Universität München
Hochschule für Politik
Political Data Science

Munich, 5. November 2020

political
data
science



Text Representations

Transform the text into a numerical form so that it can be fed into NLP and ML algorithms

A good text representation will give you a great improvement in your model! Important to choose wisely.
(Otherwise: Garbage in, garbage out)

Word classes

Simplest method to convert text into numbers.

Example:

Class1 = {Rome, Berlin, Paris, Madrid}

Class2 = {Italy, Germany, France, Spain}

One-Hot Encoding

Each word w in the corpus vocabulary (V) is given a unique integer Id between 1 and the length of the vocabulary ($|V|$)

Pet	Cat	Dog	Turtle	Fish
Cat	1	0	0	0
Dog	0	1	0	0
Turtle	0	0	1	0
Fish	0	0	0	1
Cat	1	0	0	0

One-hot encoding is the way to go if the number of words is lower than 10. For example, when a feature is **categorical**

However, normally the vocabulary is large and it will be inefficient to store, compute and to learn from (sparsity leads to overfitting)

No notion of (dis)similarity between words

Bag of Words

It begins with one-hot encoding of the words. A document is represented as a vector of size $|V|$ with counts of words as features.

Example:

Vocabulary = (Monday, Tuesday, is, a, today)

Monday Monday = [2 0 0 0 0]

Today is a Monday = [1 0 1 1 1]

Monday is today Tuesday is = [1 1 2 0 1]

Documents having the same words will have similar vector representations. Still does not capture similarity between words and the word order is lost.

Bag of N-Grams

N-Gram N continuous words (2-gram is a bigram, 3-gram is a trigram)

Example:

D1) dog bites man

D2) man bites dog

D3) dog eats meat

D4) man eats food

Vocabulary = {dog bites, bites man, man bites, bites dog, dog eats, eats meat, man eats, eats food}

D1 = [1, 1, 0, 0, 0, 0, 0, 0]

D2 = [0, 0, 1, 1, 0, 0, 0, 0]

Now order is important and captures some context. However, vocabulary size is now even larger!

Bag of words is the special case of $N=1$.

TF-IDF

Bag of words treats all words equally. **Term Frequency-Inverse Document Frequency** quantifies the importance of a word relative to other words in the document and in the corpus.

Intuition: If a word w appears many times in document D_j but does not occur much in the rest of the documents, then the word w must be of great importance to document D_j

Term Frequency How often a term occurs in a document

Inverse Document Frequency It weighs down terms that are very common accross a corpus.

Example, word w

$$TF(w, d) = \frac{\#w \text{ in } d}{\#words \text{ in } d}$$

$$IDF(w) = \frac{\#documents \text{ in corpus}}{\#documents \text{ that include word } w}$$

$$TFIDF = TF * \log(IDF)$$


Word Embeddings

Use neural network architectures to create dense, low-dimensional representations.

Intuition: Words that occur in similar contexts have similar meanings.

Example: dog and cat occur in similar contexts (I have a dog/cat as a pet)

They create a **vector space** where common vector operations lead to contextual meanings.



Word Embeddings

<i>Expression</i>	<i>Nearest token</i>
Paris - France + Italy	Rome
bigger - big + cold	colder
sushi - Japan + Germany	bratwurst
Cu - copper + gold	Au
Windows - Microsoft + Google	Android

Note: compound words like *New York* should be better pre-processed as *New_York*

How to create good word embeddings?

Word2Vec

Takes a large corpus of text as input (for example Wikipedia) and learns to represent the words based on the contexts in which they appear

Two similar variants:

- ▶ Skipgram
- ▶ CBOW

Intuition: Train a simple neural network with a single hidden layer to perform a certain task. However, the model and the outputs are not important. Instead, the goal is actually just to learn the weights of the hidden layer. These weights will be the word embeddings.

Nowadays, no state-of-the-art anymore. However, important to understand how they work. They revolutionized NLP!

Skipgram

Task: Predict the context words from the center word.

Context words defined by a **sliding window**. Normally size 5 (two to the left, and two to the right)

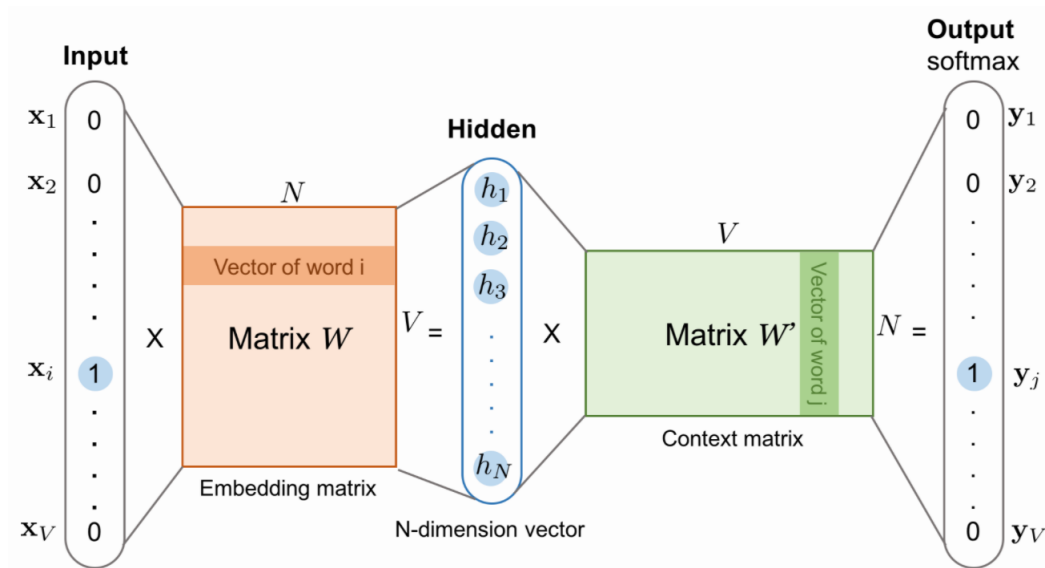
Source Text

Training Samples

Window Position	Source Text	Training Samples
1	The quick brown fox jumps over the lazy dog.	(the, quick) (the, brown)
2	The quick brown fox jumps over the lazy dog.	(quick, the) (quick, brown) (quick, fox)
3	The quick brown fox jumps over the lazy dog.	(brown, the) (brown, quick) (brown, fox) (brown, jumps)
4	The quick brown fox jumps over the lazy dog.	(fox, quick) (fox, brown) (fox, jumps) (fox, over)

Skipgram


Neural Network architecture for a vocabulary of V words and dimension of word embedding N :



Input and output are one-hot encoding vectors. W = word embedding matrix. W' = context embedding matrix

Skipgram

The Hidden Layer We are only multiplying one-hot vector to a matrix. This means we select the row of the matrix.

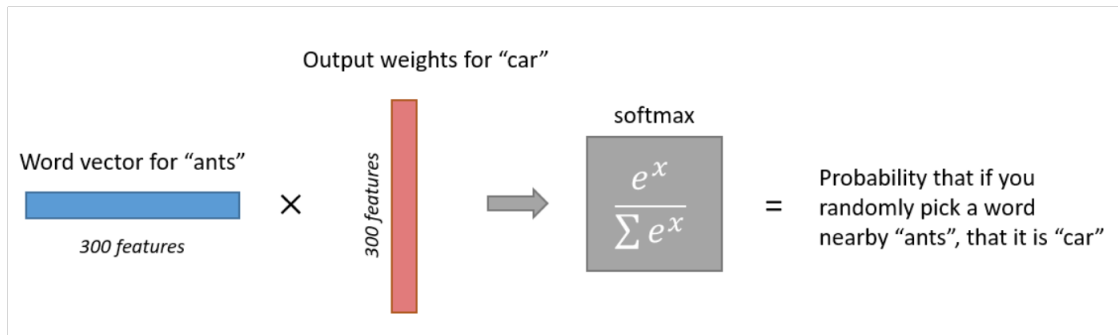


Always think of weights between neural network layers as matrices! The weight between word 800 and neuron 1 is in position $W_{800,1}$, and neuron 2 is in $W_{800,2}$

Each neuron represents one dimension of the word embedding

Skipgram

The Output Layer Each output neuron (one per word in our vocabulary!) will produce an output between 0 and 1, and the sum of all these output values will add up to 1 (using softmax)



If two different words have very similar contexts, then the model needs to output very similar results for these two words = the two words should have similar word embeddings.

Skipgram

Training this network is a hard task because of the huge number of weights.

How to make it faster and more efficient?

1. **Subsampling frequent words** Common words like "the" do not convey much context. Subsampling means that each word has a probability to be erased. The higher frequency words have a higher probability of being erased.
2. **Negative Sampling** Instead of having one output with 1 and 99,999 with 0, randomly select just a small number of “negative” words (have 0 as output) to update the weights for one training sample (for example, 5 negative words).
More frequent words are more likely to be selected as negative sample
3. **Hierarchical Softmax** Uses a tree structure to make the sum calculation faster. (Computer Science power!)

CBOW

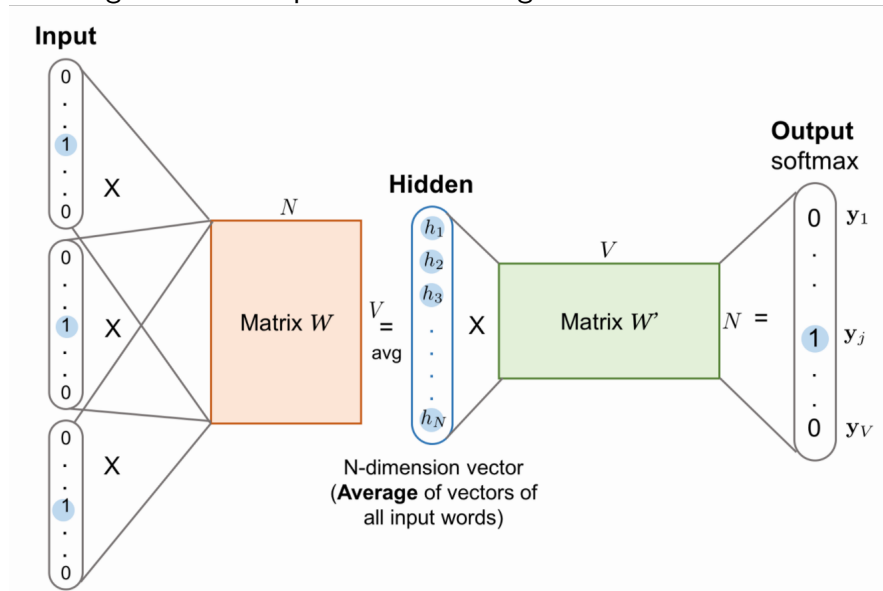
Now the other way around, use the context words as inputs and the center word as output:

Source Text	Training Samples (context, target)
The quick brown fox jumps over the lazy dog.	→ ((quick, brown), The)
The quick brown fox jumps over the lazy dog.	→ ((The, brown, fox), quick)
The quick brown fox jumps over the lazy dog.	→ ((The, quick, fox, jumps), brown)
The quick brown fox jumps over the lazy dog.	→ ((quick, brown, jumps, over), fox)

CBOW

Few changes in architecture:

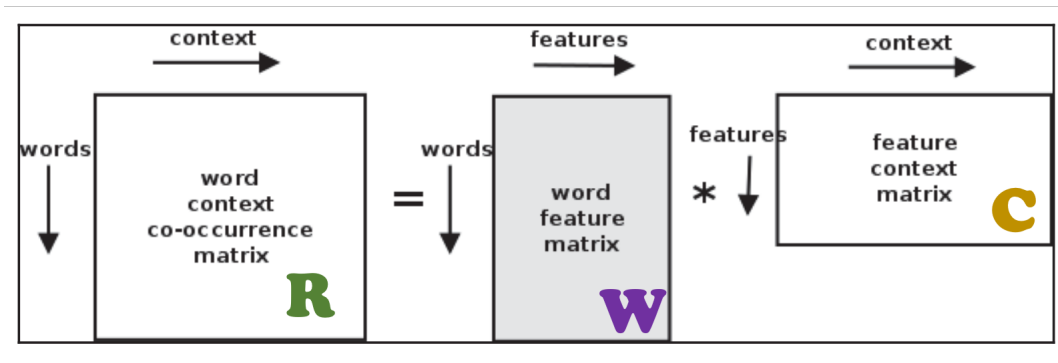
We use several one-hot encoding vectors as input and we average the word vectors of the context words



GloVe

A non-neural network approach to get word embeddings. Better results than cbow and skipgram.

Based on **matrix factorization** and **stochastic gradient descent**. Uses the **global** co-occurrences of the words.



Initialize W and C randomly. Multiply them (R') and compare with R . The difference tells us how much to change W and C to be more similar to R . Alternate between optimizing W and C .

GloVe implements a smart trick inspired by the word2vec model. It changes the **loss function** to use the ratios of co-occurrence probabilities rather than the probabilities themselves.

(What?? Go and read the original GloVe paper to get an idea of what this means)

FastText

The creator of Word2vec (Mikolov) was not happy that a non-neural approach was performing better...

He improved word2vec -> FastText

Two main tricks:

- ▶ Uses position weights. Words before and after have different weights, also closer context words than far away ones.
- ▶ Exploits **sub-word** information. (morphology)

Has better results than GloVe (but not huge improvements)

Subword Information

Representing words using its **character n-grams**. Incorporates information about the structure of the word.

Idea: same as word2vec but now the word is computed summing the ngram embeddings

Example: disastrous and disaster will have similar word embeddings because of: dis, isa, sas, ast, disa, isas, sast, di, is, sa, as, st, disas, isast

Advantage: Build vectors for unseen words. -> called Out-of-Vocabulary (OOV) words

Evaluating Word Embeddings

How do we know which embeddings are better?

Benchmarking tasks:

- ▶ **Word Analogies** List of human generated analogies (Paris -> France / Warsaw -> ?)
- ▶ **Word similarity** List of human generated similarity. Compare cosine distance between word embeddings. Normally, rare words (RW) are used to benchmark.
- ▶ **Evaluate on downstream task.** A downstream task means using the embeddings for a NLP task like classification with the same algorithm and see which embeddings give best metrics. For example, SQuAD (Stanford Question Answering Dataset)

Model	Analogy	RW	Squad
GloVe Wiki + news	72%	0.38	77.7%
fastText Wiki + news	87%	0.50	78.8%
GloVe Crawl	75%	0.52	78.9%
fastText Crawl	85%	0.58	79.8%

Visualizing Word Embeddings

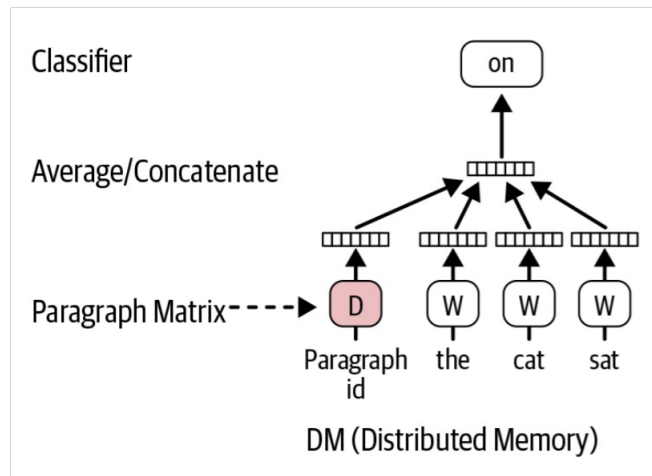
Use a **dimensionality reduction** algorithm like PCA, T-SNE or UMAP! Use first two or first three components to visualize.

`http://projector.tensorflow.org/`

Doc2Vec

Goal: Create a numeric representation of a document (a document can be a sentence, a paragraph, or a complete piece of text)

Same as CBOW, but adds a document-unique **token**. The vector representation of the token intends to represent the concept of a document



Documents with similar topics will have similar embeddings, given that they are composed of similar words.

Contextual Word Embeddings

State-of-the-art embeddings also use the context to generate embeddings.

Until now, words get one **fixed** representation. However, words can have different meanings depending on the context. (homonyms). *bat* for baseball and **bat** the animal.

Idea: Use **language modeling** to pre-train contextual embeddings.

Next week: Language Modeling

In four weeks: Transformers

Final note Hand-crafted feature representations are still important in domain-specific tasks. For example, grammar correction needs custom features to incorporate domain knowledge.